

EN3160_Assignment_01_230147L

September 13, 2026

1 EN3160 Assignment 01

Intensity Transformations and Neighborhood Filtering

Name: Dilhara D.S.

Index Number: 230147L

GitHub: [repo link](#)

```
[52]: import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt
def plt_1(a, b, x, y, xlabel, ylabel, title):
    plt.figure(figsize=(a, b)), plt.plot(x, y, '-o', markersize=3)
    plt.xlabel(xlabel), plt.ylabel(ylabel), plt.title(title)
    plt.grid(True), plt.tight_layout(), plt.show()

def plt_2(a, b, c, im, title, cmap=None):
    plt.subplot(a, b, c), plt.imshow(im, cmap=cmap)
    plt.title(title), plt.axis("off")

def plt_3(a, b, c, x, y=None, xlabel='', ylabel='', title='', kind='plot',
    ↪bins=256, range_=None):
    plt.subplot(a, b, c)
    if kind == 'hist': plt.hist(np.asarray(x).ravel(), bins=bins, range=range_)
    else: plt.plot(x, y, '-o', markersize=3)
    plt.xlabel(xlabel), plt.ylabel(ylabel)
    plt.title(title), plt.grid(True)
```

2 Question 1 — Intensity Transformation

```
[53]: def intensity_transformation(im, breakpoints):
    im_float = im.astype(np.float32)
    bp = np.asarray(breakpoints, dtype=np.float32)
    output = np.zeros_like(im_float)
    for i in range(len(bp) - 1):
        x1, y1 = bp[i]
```

```


x2, y2 = bp[i + 1]
if x2 == x1: continue
if i == len(bp) - 2: mask = (im_float >= x1) & (im_float <= x2)
else: mask = (im_float >= x1) & (im_float < x2)
output[mask] = ((y2 - y1) / (x2 - x1)) * (im_float[mask] - x1) + y1
return np.clip(output, 0, 255).astype(np.uint8)

```

```

[54]: im1 = cv.imread('images/q1.jpeg', cv.IMREAD_GRAYSCALE)
breakpoints = [(0, 0), (50, 50), (50, 100), (150, 255), (150, 150), (255, 255)]
result1 = intensity_transformation(im1, breakpoints)
x, y = [p[0] for p in breakpoints], [p[1] for p in breakpoints]
plt_1(6, 4, x, y, 'Input Intensity', 'Output Intensity', 'Intensity_
↳Transformation')

```



```

[55]: plt.figure(figsize=(6, 4))
plt_2(1, 2, 1, im1, 'Original Image', 'gray')
plt_2(1, 2, 2, result1, 'Transformed Image', 'gray')
plt.tight_layout()
plt.show()

```

Original Image




Transformed Image



3 Question 2 — White Matter and Gray Matter Enhancement


```
[56]: brain = cv.imread('images/q2.jpeg', cv.IMREAD_GRAYSCALE)
white_breakpoints = [(0, 0), (80, 20), (120, 255), (255, 255)]
white_result = intensity_transformation(brain, white_breakpoints)
gray_breakpoints = [(0, 0), (70, 1), (100, 255), (150, 255), (255, 255)]
gray_result = intensity_transformation(brain, gray_breakpoints)

x_1, y_1 = [p[0] for p in white_breakpoints], [p[1] for p in white_breakpoints]
x_2, y_2 = [p[0] for p in gray_breakpoints], [p[1] for p in gray_breakpoints]
plt.figure(figsize=(8, 4))
plt_3(1, 2, 1, x_1, y_1, 'Input Intensity', 'Output Intensity', 'White Matter_
↳Transformation')
plt_3(1, 2, 2, x_2, y_2, 'Input Intensity', 'Output Intensity', 'Gray Matter_
↳Transformation')
plt.tight_layout(), plt.show()
```



[56]: (None, None)

```
[57]: plt.figure(figsize=(6, 4))
plt_2(1, 3, 1, brain, 'Original Brain Image', 'gray')
plt_2(1, 3, 2, white_result, 'White Matter Image', 'gray')
plt_2(1, 3, 3, gray_result, 'Gray Matter Image', 'gray')
plt.tight_layout(), plt.show()
```




[57]: (None, None)

4 Question 3 — Gamma Correction


```
[58]: img3 = cv.imread('images/WhatsApp Image 2026-09-09 at 20.02.44 (1).jpeg')
lab = cv.cvtColor(img3, cv.COLOR_BGR2LAB)
L, a, b = cv.split(lab)
gamma = 0.7
L_normalized = L.astype(np.float32) / 255.0
L_gamma_corrected = np.power(L_normalized, gamma)
L_corrected = np.clip(L_gamma_corrected * 255, 0, 255).astype(np.uint8)
lab_corrected = cv.merge((L_corrected, a, b))
img3_corrected = cv.cvtColor(lab_corrected, cv.COLOR_LAB2BGR)

plt.figure(figsize=(6, 4))
plt_2(1, 2, 1, cv.cvtColor(img3, cv.COLOR_BGR2RGB), 'Original Image')
plt_2(1, 2, 2, cv.cvtColor(img3_corrected, cv.COLOR_BGR2RGB), f'Gamma Corrected_
↳Image ( = {gamma})')
plt.tight_layout(), plt.show()
```



[58]: (None, None)

```
[59]: plt.figure(figsize=(8, 4))
plt_3(1, 2, 1, L, None, 'Intensity Value', 'Frequency', 'Histogram of Original_
↳L Channel', kind='hist', bins=256, range_=(0, 256))
plt_3(1, 2, 2, L_corrected, None, 'Intensity Value', 'Frequency', 'Histogram of_
↳Gamma-Corrected L Channel', kind='hist', bins=256, range_=(0, 256))
plt.tight_layout(), plt.show()
```



[59]: (None, None)


A gamma value smaller than 1 brightens darker intensity values more strongly than brighter values. In this experiment, $\gamma = 0.7$ was used as a starting value. The corrected L histogram is shifted toward higher intensity values, indicating that darker regions have become brighter.

5 Question 4 — Vibrance Enhancement

```
[60]: img4 = cv.imread('images/WhatsApp Image 2026-09-09 at 20.02.441.jpeg')
hsv = cv.cvtColor(img4, cv.COLOR_BGR2HSV)
H, S, V = cv.split(hsv)
sigma = 70
a_value = 0.6
S_float = S.astype(np.float32)
S_new = (S_float + a_value * 128 * np.exp(-((S_float - 128) ** 2) / (2 * sigma_
    ↪** 2)))
S_new = np.clip(S_new, 0, 255).astype(np.uint8)
hsv_new = cv.merge((H, S_new, V))
vibrance = cv.cvtColor(hsv_new, cv.COLOR_HSV2BGR)
x = np.arange(256)
transformation = np.minimum(255, x + a_value * 128 * np.exp(-((x - 128) ** 2) /
    ↪(2 * sigma ** 2)))

plt_1(4, 4, x, transformation, 'Input Intensity', 'Output Intensity',
    ↪f'Vibrance Transformation Curve (a = {a_value}, sigma = {sigma})')
plt.figure(figsize=(6, 4))
plt_2(1, 2, 1, cv.cvtColor(img4, cv.COLOR_BGR2RGB), 'Original Image')
plt_2(1, 2, 2, cv.cvtColor(vibrance, cv.COLOR_BGR2RGB), f'Vibrance Adjusted_
    ↪Image (a = {a_value}, sigma = {sigma})')
```

```
plt.tight_layout(), plt.show()
```



[60]: (None, None)

The saturation plane is increased mainly around the middle of the saturation range. The hue and value planes remain unchanged.

A value of $a = 0.6$ was selected as an initial visually pleasing value. Increasing a produces stronger color enhancement, while a smaller value produces a more subtle effect.


The final value of a should be selected based on the actual appearance of the supplied image.

6 Question 5 — Histogram Equalization

Histogram equalization is implemented manually.


```
[61]: def histogram_equalization(im):  
    hist = np.bincount(im.ravel(), minlength=256)  
    cdf = np.cumsum(hist)  
    nonzero = cdf[cdf > 0]  
    if len(nonzero) == 0: return im.copy()  
    cdf_min = nonzero[0]  
    total_pixels = im.size  
    denominator = total_pixels - cdf_min  
    if denominator == 0: return im.copy()  
    lut = ((cdf - cdf_min) / denominator * 255)  
    lut = np.clip(lut, 0, 255).astype(np.uint8)  
    equalized_im = lut[im]  
    return equalized_im
```

```
[62]: img5 = cv.imread('images/shells.tif', cv.IMREAD_GRAYSCALE)  
equalized_img5 = histogram_equalization(img5)  
plt.figure(figsize=(6, 4))  
plt_2(1, 2, 1, img5, 'Original Image', 'gray')  
plt_2(1, 2, 2, equalized_img5, 'Equalized Image', 'gray')  
plt.tight_layout(), plt.show()
```



```
[62]: (None, None)
```

```
[63]: plt.figure(figsize=(8, 4))
plt_3(1, 2, 1, img5, None, 'Intensity Value', 'Frequency', 'Histogram of_
↳Original Image', kind='hist', bins=256, range_=(0, 256))
plt_3(1, 2, 2, equalized_img5, None, 'Intensity Value', 'Frequency', 'Histogram_
↳of Equalized Image', kind='hist', bins=256, range_=(0, 256))
plt.tight_layout(), plt.show()
```



```
[63]: (None, None)
```

7 Question 6 — Histogram Equalization of the Foreground

In this question, histogram equalization is applied only to the foreground.

```
[64]: img6 = cv.imread('images/WhatsApp Image 2026-09-09 at 20.02.44 (2).jpeg')
hsv_img6 = cv.cvtColor(img6, cv.COLOR_BGR2HSV)
H6, S6, V6 = cv.split(hsv_img6)
plt.figure(figsize=(15, 4))
plt_2(1, 3, 1, H6, 'Hue Plane', 'gray')
plt_2(1, 3, 2, S6, 'Saturation Plane', 'gray')
plt_2(1, 3, 3, V6, 'Value Plane', 'gray')
plt.tight_layout(), plt.show()
```



[64]: (None, None)

```
[65]: threshold_value = 80
_, mask6 = cv.threshold(S6, threshold_value, 255, cv.THRESH_BINARY)
kernel = np.ones((5, 5), np.uint8)
mask6 = cv.morphologyEx(mask6, cv.MORPH_OPEN, kernel)
mask6 = cv.morphologyEx(mask6, cv.MORPH_CLOSE, kernel)
plt.figure(figsize=(5, 4))
plt_2(1, 1, 1, mask6, 'Foreground Mask', 'gray')
plt.tight_layout(), plt.show()
```

Foreground Mask



[65]: (None, None)

```
[66]: foreground6 = cv.bitwise_and(img6, img6, mask=mask6)
gray6 = cv.cvtColor(foreground6, cv.COLOR_BGR2GRAY)
foreground_pixels = gray6[mask6 > 0]
hist6 = np.bincount(foreground_pixels, minlength=256)
cdf6 = np.cumsum(hist6)
nonzero = cdf6[cdf6 > 0]
if len(nonzero) == 0:
    raise ValueError("No foreground pixels were detected. Adjust the threshold.
↪")
cdf_min = nonzero[0]
total_foreground_pixels = foreground_pixels.size
denominator = total_foreground_pixels - cdf_min
```

```


if denominator == 0:
    lut6 = np.arange(256, dtype=np.uint8)
else:
    lut6 = ((cdf6 - cdf_min) / denominator * 255)
    lut6 = np.clip(lut6, 0, 255).astype(np.uint8)
equalized_foreground_gray = np.zeros_like(gray6)
equalized_foreground_gray[mask6 > 0] = lut6[gray6[mask6 > 0]]
background_mask6 = cv.bitwise_not(mask6)
background6 = cv.bitwise_and(img6, img6, mask=background_mask6)
equalized_foreground6 = cv.cvtColor(equalized_foreground_gray, cv.
    ↪COLOR_GRAY2BGR)
result6 = background6.copy()
result6[mask6 > 0] = equalized_foreground6[mask6 > 0]

```

```

[67]: plt.figure(figsize=(15, 8))
plt_2(2, 4, 1, H6, 'Hue', 'gray')
plt_2(2, 4, 2, S6, 'Saturation', 'gray')
plt_2(2, 4, 3, V6, 'Value', 'gray')
plt_2(2, 4, 4, mask6, 'Foreground Mask', 'gray')
plt_2(2, 4, 5, cv.cvtColor(img6, cv.COLOR_BGR2RGB), 'Original Image')
plt_2(2, 4, 6, cv.cvtColor(foreground6, cv.COLOR_BGR2RGB), 'Foreground')
plt_2(2, 4, 7, cv.cvtColor(background6, cv.COLOR_BGR2RGB), 'Background')
plt_2(2, 4, 8, cv.cvtColor(result6, cv.COLOR_BGR2RGB), 'Equalized Foreground')
plt.tight_layout(), plt.show()

```



[67]: (None, None)

8 Question 7 — Sobel Filtering

```
[68]: img7 = cv.imread('images/WhatsApp Image 2026-09-09 at 20.02.43 (2).jpeg', cv.
      ↪IMREAD_GRAYSCALE)
sobel_x_kernel = np.array([[ -1,  0,  1],
                           [ -2,  0,  2],
                           [ -1,  0,  1]], dtype=np.float32)
sobel_y_kernel = np.array([[ -1, -2, -1],
                           [  0,  0,  0],
                           [  1,  2,  1]], dtype=np.float32)
gx = cv.filter2D(img7, cv.CV_32F, sobel_x_kernel)
gy = cv.filter2D(img7, cv.CV_32F, sobel_y_kernel)
magnitude = np.sqrt(gx**2 + gy**2)
magnitude = cv.normalize(magnitude, None, 0, 255, cv.NORM_MINMAX).astype(np.
      ↪uint8)
plt.figure(figsize=(12, 4))
plt_2(1, 4, 1, img7, 'Original Image', 'gray')
plt_2(1, 4, 2, cv.convertScaleAbs(gx), 'Sobel X', 'gray')
plt_2(1, 4, 3, cv.convertScaleAbs(gy), 'Sobel Y', 'gray')
plt_2(1, 4, 4, magnitude, 'Gradient Magnitude', 'gray')
plt.tight_layout(), plt.show()
```



[68]: (None, None)

```
[69]: def filter2D_custom(im, kernel):
      im = im.astype(np.float32)
      kernel = kernel.astype(np.float32)
      kh, kw = kernel.shape
      pad_h, pad_w = kh // 2, kw // 2
      padded_im = np.pad(im, ((pad_h, pad_h), (pad_w, pad_w)), mode='edge')
      h, w = im.shape
      result = np.zeros((h, w), dtype=np.float32)
      for i in range(h):
          for j in range(w):
              region = padded_im[i:i + kh, j:j + kw]
```

```

        result[i, j] = np.sum(region * kernel)
    return result

```

```


[70]: gx_custom = filter2D_custom(img7, sobel_x_kernel)
      gy_custom = filter2D_custom(img7, sobel_y_kernel)
      magnitude_custom = np.sqrt(gx_custom**2 + gy_custom**2)
      magnitude_custom = cv.normalize(magnitude_custom, None, 0, 255, cv.NORM_MINMAX).
      ↪astype(np.uint8)
      plt.figure(figsize=(10, 4))
      plt_2(1, 2, 1, magnitude, 'OpenCV Sobel Magnitude', 'gray')
      plt_2(1, 2, 2, magnitude_custom, 'Custom Sobel Magnitude', 'gray')
      plt.tight_layout(), plt.show()

```

OpenCV Sobel Magnitude



Custom Sobel Magnitude



```

[70]: (None, None)

```

```

[71]: sobel_mse = np.mean((magnitude.astype(np.float32) - magnitude_custom.astype(np.
      ↪float32)) ** 2)
      print(f"MSE between OpenCV and custom Sobel magnitude: {sobel_mse:.6f}")

```

MSE between OpenCV and custom Sobel magnitude: 0.592199

8.1 7(c) Separable Sobel Filtering

```


[72]: smooth_vertical = np.array([[1], [2], [1]], dtype=np.float32)
      derivative_horizontal = np.array([[1, 0, -1]], dtype=np.float32)
      temp_x = cv.filter2D(img7, cv.CV_32F, smooth_vertical)
      gx_separable = cv.filter2D(temp_x, cv.CV_32F, derivative_horizontal)
      smooth_horizontal = np.array([[1, 2, 1]], dtype=np.float32)
      derivative_vertical = np.array([[1], [0], [-1]], dtype=np.float32)

```

```

temp_y = cv.filter2D(img7, cv.CV_32F, smooth_horizontal)
gy_separable = cv.filter2D(temp_y, cv.CV_32F, derivative_vertical)
magnitude_separable = np.sqrt(gx_separable**2 + gy_separable**2)
magnitude_separable = cv.normalize(magnitude_separable, None, 0, 255, cv.
    ↪ NORM_MINMAX).astype(np.uint8)
plt.figure(figsize=(12, 4))
plt_2(1, 3, 1, magnitude, '2-D Sobel Magnitude', 'gray')
plt_2(1, 3, 2, magnitude_custom, 'Custom 2-D Magnitude', 'gray')
plt_2(1, 3, 3, magnitude_separable, 'Separable Sobel Magnitude', 'gray')
plt.tight_layout(), plt.show()

```



[72]: (None, None)

9 Question 8 — Intensity Transformation

```

[73]: def zoom_nearest(im, s):
    if not (0 < s <= 10): raise ValueError("Scale s must satisfy 0 < s <= 10")
    h, w = im.shape[:2]
    new_h, new_w = int(round(h * s)), int(round(w * s))
    y_idx = np.minimum((np.arange(new_h) / s).astype(int), h - 1)
    x_idx = np.minimum((np.arange(new_w) / s).astype(int), w - 1)
    return im[y_idx[:, None], x_idx[None, :]]

def zoom_bilinear(im, s):
    if not (0 < s <= 10): raise ValueError("Scale s must satisfy 0 < s <= 10")
    if im.ndim != 2: raise ValueError("zoom_bilinear expects a grayscale image")
    h, w = im.shape
    new_h, new_w = int(round(h * s)), int(round(w * s))
    output = np.zeros((new_h, new_w), dtype=np.float32)
    for y in range(new_h):
        src_y = y / s
        y0 = min(int(np.floor(src_y)), h - 1)

```

```

        y1 = min(y0 + 1, h - 1)
        dy = src_y - y0
        for x in range(new_w):
            src_x = x / s
            x0 = min(int(np.floor(src_x)), w - 1)
            x1 = min(x0 + 1, w - 1)
            dx = src_x - x0
            top = (1 - dx) * im[y0, x0] + dx * im[y0, x1]
            bottom = (1 - dx) * im[y1, x0] + dx * im[y1, x1]
            output[y, x] = (1 - dy) * top + dy * bottom
        return np.clip(output, 0, 255).astype(im.dtype)

def normalized_ssd(original, reconstructed):
    original = original.astype(np.float32)
    reconstructed = reconstructed.astype(np.float32)
    return np.mean((original - reconstructed) ** 2) / (255.0 ** 2)

```


```

[74]: scale = 4
zoom_results = []
for number in (1, 2, 3):
    small_path = f'images/im0{number}small.png'
    large_path = f'images/im0{number}.png'
    small = cv.imread(small_path, cv.IMREAD_GRAYSCALE)
    large = cv.imread(large_path, cv.IMREAD_GRAYSCALE)
    if small is None or large is None:
        print(f"Skipping pair {number}: check {small_path} and {large_path}")
        continue
    nearest = zoom_nearest(small, scale)
    bilinear = zoom_bilinear(small, scale)
    if nearest.shape != large.shape:
        nearest = cv.resize(nearest, (large.shape[1], large.shape[0]),
↪ interpolation=cv.INTER_NEAREST)
    if bilinear.shape != large.shape:
        bilinear = cv.resize(bilinear, (large.shape[1], large.shape[0]),
↪ interpolation=cv.INTER_LINEAR)
    ssd_nearest = normalized_ssd(large, nearest)
    ssd_bilinear = normalized_ssd(large, bilinear)
    zoom_results.append((number, large, nearest, bilinear, ssd_nearest,
↪ ssd_bilinear))
    print(f"Image pair {number}: nearest SSD = {ssd_nearest:.6f}, bilinear SSD
↪ = {ssd_bilinear:.6f}")

```

Image pair 1: nearest SSD = 0.002110, bilinear SSD = 0.003098
Image pair 2: nearest SSD = 0.000410, bilinear SSD = 0.000762
Image pair 3: nearest SSD = 0.001027, bilinear SSD = 0.001707

```
[75]: for number, large, nearest, bilinear, ssd_nearest, ssd_bilinear in zoom_results:
plt.figure(figsize=(15, 4))
plt_2(1, 3, 1, large, f'Original Large Image {number}', 'gray')
plt_2(1, 3, 2, nearest, f'Nearest Neighbor\nSSD = {ssd_nearest:.6f}', 'gray')
plt_2(1, 3, 3, bilinear, f'Bilinear\nSSD = {ssd_bilinear:.6f}', 'gray')
plt.tight_layout(), plt.show()
```



10 Question 9 — GrabCut Segmentation

```
[76]: img9 = cv.imread(r"images\WhatsApp Image 2026-09-09 at 20.02.43 (1).jpeg")
h, w = img9.shape[:2]
mask9 = np.zeros((h, w), np.uint8)
bgModel = np.zeros((1, 65), np.float64)
fgModel = np.zeros((1, 65), np.float64)
```

```


rect = (int(w * 0.10), int(h * 0.05), int(w * 0.80), int(h * 0.85))
cv.grabCut(img9, mask9, rect, bgModel, fgModel, 5, cv.GC_INIT_WITH_RECT)
binary_mask9 = np.where((mask9 == cv.GC_FGD) | (mask9 == cv.GC_PR_FGD), 255, 0).
    ↳astype(np.uint8)
foreground9 = cv.bitwise_and(img9, img9, mask=binary_mask9)
background_mask9 = cv.bitwise_not(binary_mask9)
background9 = cv.bitwise_and(img9, img9, mask=background_mask9)

```

```

[77]: plt.figure(figsize=(15, 4))
plt_2(1, 3, 1, binary_mask9, 'GrabCut Mask', 'gray')
plt_2(1, 3, 2, cv.cvtColor(foreground9, cv.COLOR_BGR2RGB), 'Foreground')
plt_2(1, 3, 3, cv.cvtColor(background9, cv.COLOR_BGR2RGB), 'Background')
plt.tight_layout(), plt.show()

```



[77]: (None, None)

```

[78]: blurred_background9 = cv.GaussianBlur(img9, (51, 51), 0)
enhanced9 = np.where(binary_mask9[:, :, None] == 255, img9, blurred_background9)
enhanced9 = enhanced9.astype(np.uint8)

```

```

[79]: plt.figure(figsize=(10, 4))
plt_2(1, 2, 1, cv.cvtColor(img9, cv.COLOR_BGR2RGB), 'Original Image')
plt_2(1, 2, 2, cv.cvtColor(enhanced9, cv.COLOR_BGR2RGB), 'Enhanced Image')
plt.tight_layout(), plt.show()

```

Original Image



Enhanced Image



[79]: (None, None)

10.1 Why is the background just beyond the flower edge dark?

The dark region occurs because the GrabCut segmentation mask does not perfectly follow the boundary of the flower.

Some pixels immediately outside the flower may be classified as background. These pixels are replaced by the blurred background.

Since the surrounding background contains dark regions, Gaussian blurring spreads the intensity of these regions around the boundary.


Therefore, the dark area is mainly caused by segmentation boundary errors combined with strong background blurring.

11 Question 10 — Intensity Transformation

```
[80]: img10 = cv.imread(r"images\WhatsApp Image 2026-09-09 at 20.02.44 (3).jpeg")
      bilateral10 = cv.bilateralFilter(img10, d=9, sigmaColor=75, sigmaSpace=75)
      gaussian10 = cv.GaussianBlur(img10, (9, 9), 75)
```

```
[81]: plt.figure(figsize=(15, 4))
      plt_2(1, 3, 1, cv.cvtColor(img10, cv.COLOR_BGR2RGB), 'Original Image')
      plt_2(1, 3, 2, cv.cvtColor(gaussian10, cv.COLOR_BGR2RGB), 'Gaussian Blur')
      plt_2(1, 3, 3, cv.cvtColor(bilateral10, cv.COLOR_BGR2RGB), 'OpenCV Bilateral_
      ↪Filter')
```

```
plt.tight_layout(), plt.show()
```



```
[81]: (None, None)
```


11.1 10(b) Own Bilateral Filter

```
[82]: def bilateral_filter_gray(image, kernel_size=5, sigma_space=10, sigma_range=30):
    image = image.astype(np.float32)
    pad = kernel_size // 2
    padded = np.pad(image, pad, mode="reflect")
    h, w = image.shape
    output = np.zeros_like(image)
    ax = np.arange(-pad, pad + 1)
    xx, yy = np.meshgrid(ax, ax)
    spatial_weights = np.exp(-(xx ** 2 + yy ** 2) / (2 * sigma_space ** 2))
    for i in range(h):
        for j in range(w):
            region = padded[i:i + kernel_size, j:j + kernel_size]
            center = image[i, j]
            range_weights = np.exp(-((region - center) ** 2) / (2 * sigma_range
↪** 2))

            weights = (spatial_weights * range_weights)
            weighted_sum = np.sum(weights * region)
            weight_sum = np.sum(weights)
            output[i, j] = (weighted_sum / weight_sum)
    return np.clip(output, 0, 255).astype(np.uint8)
```

```
[83]: gray10 = cv.cvtColor(img10, cv.COLOR_BGR2GRAY)
opencv_bilateral_gray = cv.bilateralFilter(gray10, d=5, sigmaColor=30,
↪sigmaSpace=10)
my_bilateral_gray = bilateral_filter_gray(gray10, kernel_size=5,
↪sigma_space=10, sigma_range=30)
```

```
plt.figure(figsize=(15, 4))
plt_2(1, 3, 1, gray10, 'Original Grayscale', 'gray')
plt_2(1, 3, 2, opencv_bilateral_gray, 'OpenCV Bilateral', 'gray')
plt_2(1, 3, 3, my_bilateral_gray, 'Custom Bilateral Filter', 'gray')
plt.tight_layout(), plt.show()
```



[83]: (None, None)

11.2 Quantitative Comparison

Mean Squared Error (MSE) is used to compare the manually implemented bilateral filter with OpenCV's bilateral filter.

A smaller MSE indicates that the manually implemented result is closer to OpenCV's result.

```
[84]: mse10 = np.mean((opencv_bilateral_gray.astype(np.float32) - my_bilateral_gray.
    ↳astype(np.float32)) ** 2)
print("MSE between OpenCV and my bilateral filter:", mse10)
if mse10 == 0: psnr10 = float("inf")
else: psnr10 = 10 * np.log10((255 ** 2) / mse10)
print("PSNR:", psnr10, "dB")
```

MSE between OpenCV and my bilateral filter: 5.617387
PSNR: 40.63546 dB